AdSafe Compliance Policy Framework

Ad Compliance & Brand Safety for Social Video Ads

Policy Reference Document · Version 1.0 · February 2026

Overview

This document defines the six compliance policy categories evaluated by the AdSafe automated video review system. Each policy is grounded in published industry standards and applied at a high-severity threshold: any detectable infringement in a paid advertisement is flagged, regardless of duration or apparent intent.

⚖️ Regulatory Grounding: Policies are derived from the GARM Brand Safety Floor + Suitability Framework (2022), FTC Guides Concerning Use of Endorsements and Testimonials in Advertising (revised July 2023, 16 CFR Part 255), and FDA FD&C Act cosmetic/drug claim distinctions (21 U.S.C. §321). GARM was discontinued in August 2024; its framework remains the industry standard adopted by major platforms.

Verdict Logic

The system returns one of three verdicts based on the lowest-passing policy across all categories:

Verdict	Trigger condition	Action
🚫 **BLOCK**	Any policy = fail OR off_brief	Ad is rejected. Creator notified with timestamped evidence.
⚠️ **REVIEW**	Any policy = warn OR borderline OR low confidence	Routed to human reviewer queue with AI-generated summary.
✅ **APPROVE**	All policies = pass AND on_brief AND confidence ≥ medium	Ad cleared for promotion automatically.

Policy 1 — Hate / Harassment

📋 *Source: GARM Brand Safety Floor + Suitability Framework (2022), Category 6: Hate Speech & Acts of Aggression*

Covers dehumanizing, demeaning, or hostile content directed at individuals or groups based on protected characteristics. Applied broadly to cover speech, on-screen text, visible background content, and implied meaning — not just explicit slurs.

FAIL	• Slurs, dehumanizing language, or epithets targeting race, ethnicity, religion, gender, sexual orientation, disability, or nationality — spoken, on-screen, or visible in background • Content that mocks, demeans, or stereotypes a group of people • Bullying or targeted harassment of an individual by name or appearance • Gestures, symbols, or imagery associated with hate movements • Language inciting hostility or violence toward any group • Derogatory body commentary: fat-shaming, colorism, or ageism presented approvingly
WARN	• Edgy humor that could reasonably be interpreted as demeaning, even if unintentional • Comparative language implying one skin type, tone, or feature is superior • Exclusionary framing that marginalizes a group (e.g., 'only for people who...')
PASS	• No language, imagery, or content that demeans or targets any group • Inclusive or neutral language throughout • Diversity shown respectfully without comment

Policy 2 — Profanity / Explicit Language

📋 ***Source: GARM Brand Safety Floor + Suitability Framework (2022), Category 3: Offensive Language & Profanity***

Paid advertisements representing a brand must meet broadcast-equivalent standards. Any profanity — even bleeped or mild — reflects on the brand and is flagged given the high-severity threshold applied to this campaign type.

FAIL	• Strong profanity: f-word, s-word, c-word, or equivalent in any language • Slurs used as profanity regardless of reclamation context • Sexually explicit language or graphic descriptions of sexual acts • Explicit violent speech or graphic threats

WARN	• Mild profanity: 'damn', 'hell', 'ass', 'crap', 'bitch', 'bastard' — even used casually or in passing • Bleeped or censored profanity (audible intent is still flagged) • Suggestive sexual language that stops short of explicit • Aggressive or crude language not directed at a group
PASS	• All language is clean and appropriate for a general audience • No profanity, censored or uncensored, in speech or on-screen text

Policy 3 — Drugs / Illegal Behavior

📋 *Source: GARM Brand Safety Floor + Suitability Framework (2022), Category 4: Drugs, Tobacco, e-Cigarettes, Vaping & Alcohol*

Covers all controlled substances, tobacco products, vaping, and illegal activity. Given this is a cosmetics campaign targeting general audiences, alcohol is also flagged at the WARN level when prominently featured on camera.

FAIL	• Illegal drug use depicted, referenced positively, or glorified (cannabis, cocaine, prescription misuse, etc.) • Drug paraphernalia visible: pipes, syringes, rolling papers, etc. • Tobacco or cigarette smoking shown or referenced approvingly • Vaping or e-cigarettes shown in any context • Illegal activity depicted, encouraged, or normalized • Creator appears visibly intoxicated during product demonstration
WARN	• Alcohol prominently consumed or referenced on camera (background glass = pass; drinking on camera = warn) • Prescription medication mentioned in a way that could encourage misuse • Casual, non-glorifying references to past drug or alcohol use • Content filmed in an environment where off-camera drug use is implied
PASS	• No substances, paraphernalia, or illegal activity visible or referenced • Alcohol present incidentally in background without reference or discussion

Policy 4 — Unsafe / Misleading Product Usage

📋 ***Source: Advertiser brand safety requirements; FDA cosmetic safety standards; general duty of care for advertised products***

Specific to cosmetics and beauty products. Covers incorrect, unsanitary, or dangerous application techniques that could cause consumer harm or create brand liability. Pegasus evaluates both what the creator does and what they instruct viewers to do.

FAIL	• Product applied directly to eye mucosa, waterline, or inner eyelid unless formulated and labeled for that specific use • Face product applied to the body in a way that could cause irritation or harm (e.g., face exfoliant on broken skin) • Product ingested, inhaled, or used entirely outside its intended purpose • Incompatible products layered dangerously (e.g., high-concentration AHA + retinol + undiluted essential oils presented as safe) • Unsanitary application: double-dipping in shared pots with bare hands, using visibly dirty tools, applying tester products • Undisclosed DIY modifications to a product (adding ingredients, diluting) presented as safe or recommended • Physical harm risk: aggressive skin pulling, unsafe waxing, chemical peel without proper instruction
WARN	• Application technique noticeably sloppy or inconsistent with safe usage, even if not immediately dangerous • Tools used appear unclean or could reasonably spread bacteria • Product applied adjacent to eyes without appropriate caution (e.g., serum not rated for eye area used near socket) • Combination shown that dermatologists commonly advise against without any caveat
PASS	• Product applied correctly, cleanly, and in accordance with its intended use • Creator demonstrates safe, hygienic technique throughout • No risky combinations or off-label application methods shown

Policy 5 — Medical / Cosmetic Claims

📋 ***Source: FDA FD&C Act §201(g)/(i); FTC Endorsement Guides (2023 revision, 16 CFR Part 255); FTC Act §5 (deceptive practices)***

Under the FDA FD&C Act, a cosmetic may only claim to cleanse, beautify, promote attractiveness, or alter appearance. Claims that a product treats, cures, prevents, or affects the structure or function of the body convert it into an unapproved drug — illegal to market without FDA approval. The FTC additionally prohibits unsubstantiated objective claims in advertising.

Key distinction: 'reduces the appearance of fine lines' (cosmetic claim — legal) vs. 'removes wrinkles' (structural claim — illegal drug claim). 'Leaves skin looking brighter' (cosmetic) vs. 'increases melanin production' (physiological — illegal). When in doubt, flag it.

FAIL	• Disease treatment claims: 'treats/cures/heals acne, eczema, psoriasis, rosacea, dermatitis, hyperpigmentation' • Structural/physiological claims: 'regenerates skin cells', 'stimulates collagen production', 'repairs skin barrier', 'reverses aging at the cellular level' • Physiological effect claims: 'increases/decreases melanin', 'reduces/eliminates cellulite', 'restores hair growth' • Unsubstantiated absolute claims: 'clinically proven to [outcome]' without cited study, '#1 [category]' without source, 'guaranteed to work', '100% effective' • Unverified authority claims: 'dermatologist approved/recommended' without disclosed basis • FTC disclosure failure: paid/gifted relationship not clearly and conspicuously disclosed at the START of the video • Disclosure buried in description, end-of-video, or hashtags only — or relying solely on platform disclosure tool
WARN	• Borderline structure-function language: 'helps repair skin', 'rebuilds moisture barrier', 'boosts skin's natural renewal' • 'Clinically tested' or 'dermatologist tested' without clarifying what was tested or the outcome • Exaggerated before/after comparisons without disclosure of lighting, makeup, or editing differences • Creator implies product resolved a named condition without explicitly claiming to treat it • Disclosure present but delayed (appears after 30+ seconds of content)
PASS	• All claims are appearance-based and cosmetically appropriate: 'looks smoother', 'feels hydrated', 'brightens complexion', 'reduces the appearance of fine lines' • Paid/gifted relationship disclosed clearly and early in the video • No comparative or absolute efficacy claims without substantiation • Creator shares personal experience without generalizing to all users

Policy 6 — Campaign Relevance (On-Brief Check)

📋 *Source: Advertiser campaign brief requirements; platform ad promotion standards*

Ensures that creator content being promoted as a paid ad is actually about the advertised product. Off-brief content wastes advertiser spend and misleads consumers. The threshold is deliberately strict: the product must be the clear focus, not just an incidental mention.

FAIL	• Promoted product is never shown, named, or described in the video • Video is primarily about an unrelated topic — product appears in < 15% of content • Content is clickbait: title/hook is unrelated to the product • A competitor product receives equal or more screen time than the promoted product
WARN	• Product is present but not the clear focus (15–40% of content) • Creator mentions product but spends most time on unrelated topics (haul with many products) • Campaign brief calls for tutorial but video is primarily a lifestyle vlog
PASS	• Product is clearly the subject of the video (> 40% of content focused on the product) • Content aligns with the campaign brief in tone, format, and subject matter • Creator demonstrates, reviews, or meaningfully discusses the product

Confidence Field

Every policy evaluation includes a confidence level. This is critical for the REVIEW escalation path:

HIGH	Unambiguous finding. Direct quote, clear visual, or unmistakable on-screen text supports the call. No reasonable alternative interpretation.
MEDIUM	Reasonably confident. Finding is supported by context but some ambiguity exists. Human review recommended for BLOCK verdicts.
LOW	Uncertain. Content was ambiguous, brief, or partially obscured. Automatically triggers REVIEW regardless of status. A false positive costs one review; a false negative costs brand reputation.

Why Outputs Are Trustworthy

Three mechanisms ensure the system's outputs are auditable and defensible:

1. Regulatory Grounding

Every FAIL and WARN criterion maps to a specific published standard — GARM category, FTC regulation, or FDA statute. Decisions can be explained to advertisers and legal teams by citing the exact rule violated, not just 'the AI said so.'

2. Evidence + Reasoning Fields

The structured JSON output requires both an evidence field (direct quote, visual description, or timestamp) and a reasoning field (one sentence explanation) for every policy category. This creates an audit trail — reviewers can verify the AI's call against the actual video moment.

3. Conservative Escalation

Any low-confidence finding triggers REVIEW rather than APPROVE. The system is designed to over-flag rather than under-flag: a false positive costs one human review; a false negative costs brand reputation, legal exposure, and advertiser trust. When in doubt, escalate.

Production Scaling Architecture

In a production ads pipeline, this compliance system would operate as follows:

- Creator submits video for promotion → event published to job queue (e.g., AWS SQS)
- Worker pool calls TwelveLabs indexing API; video is indexed with both Pegasus 1.2 (analysis) and Marengo 3.0 (search)
- On indexing completion webhook: Pegasus Analyze API called with full policy prompt → structured JSON verdict stored to database
- For WARN/BLOCK verdicts: Marengo Search API called per flagged policy to retrieve timestamped evidence clips
- APPROVE: ad cleared automatically, no human review
- REVIEW: routed to human reviewer queue with AI summary and timestamped clips pre-loaded
- BLOCK: creator notified with verdict, policy category, evidence quote, and timestamp — no manual review required
- Index reuse: one shared index per campaign reduces cost; videos remain searchable for audit